

Labwork 3 – Report

Member:

Phan Manh Tung – USTHBI8-160

Vu Tuan Phong – USTHBI8-139

T2 – Classification I

Data: For this labwork, we choose 2 dataset from UCI Machine Learning Repository namely Iris dataset and Breast Cancer Wisconsin dataset.

Tool: Python is our chosen programming language and we opt for google colab (a free python interactive environment offered by Google) for strong computing power and convenient coding.

I) Perceptron

1.1 Iris dataset:

In this problem, we attempt to implement the model from scratch.

First, we prepare the demanding input for the model. Iris dataset has 3 classes, which is not suitable for our binary model. Therefore, we eliminate class Iris-setosa, leaving the dataset with 100 sample divided into 2 classes namely, Iris-virginica, Iris-versicolor.


```
[181] import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

column = ["sepal length", "sepal width", "petal length", "petal width", "class"]
df = pd.read_csv("/content/gdrive/My Drive/iris.data", sep=',', header=None)
df.columns = column
myclass = df["class"]
df.head(5)
```

	sepal length	sepal width	petal length	petal width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
[183] # Get rid of class "Iris-setosa":
mydf = df[df['class'] != "Iris-setosa"]
twoclass = mydf["class"]
print(len(twoclass))
```

100



A linear classifier has the form:

$$f(\mathbf{x}) = \mathbf{w}\mathbf{x}^T + b$$

Because our activation function of Perceptron algorithm is the sign of $f(\mathbf{x})$ – the hyperplane, so that we also have to change the labels of the dataset into values 1 and -1 (-1 is Iris-versicolor and 1 is Iris-virginica).

Then we divide the dataset into train/test set with the ratio of 70:30.

```
[184] X = mydf.drop(["class"], axis=1).values
      y = twoclass.values

      # Turn the class into -1; 1 for training purpose
      # Iris-versicolor = -1; Iris-virginica = 1
      y = np.where(y == "Iris-versicolor", -1, 1)

      # splitting the data into training and test sets (70:30)
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=4)
```

Below part is our main implementation of the Perceptron algorithm. We follow the rule given on the slide.

Perceptron Algorithm

- initialize $\mathbf{w}_0 = \mathbf{0}$ (or close to $\mathbf{0}$)
- for all $\mathbf{x}_i \in \mathcal{X}$, if \mathbf{x}_i is misclassified

$$\mathbf{w}_t = \mathbf{w}_{t-1} + \alpha \text{sign}(f(\mathbf{x}_i))\mathbf{x}_i$$

- until all the data is correctly classified

```
[185] # Perceptron in action:
def h(sl, sw, pl, pw, w1, w2, w3, w4, b):
    return np.sign(sl*w1 + sw*w2 + pl*w3 + pw*w4 + b)

w1 = 0
w2 = 0
w3 = 0
w4 = 0
b = 0

alpha = 0.1
for j in range(100):
    for i in range(0, len(X_train)):
        if h(X_train[i][0], X_train[i][1], X_train[i][2], X_train[i][3], w1, w2, w3, w4, b) != y_train[i]:
            w1 = w1 + alpha*y_train[i]*X_train[i][0]
            w2 = w2 + alpha*y_train[i]*X_train[i][1]
            w3 = w3 + alpha*y_train[i]*X_train[i][2]
            w4 = w4 + alpha*y_train[i]*X_train[i][3]
            b = b + alpha*y_train[i]

print(w1)
print(w2)
print(w3)
print(w4)
print(b)
```

```
↳ -5.679999999999973
   -3.310000000000018
   6.399999999999917
   10.929999999999943
   -6.299999999999994
```

The $h()$ function is our activation function, which returns the sign of the hyperplane (value 1 or -1).

Because Iris dataset has 4 attributes: sepal length, sepal width, petal length, petal width. The hyperplane $f(x)$ will contain 4 weights $w1$, $w2$, $w3$, $w4$ and bias value b . We initialize all 5 values at 0.

The learning rate (α) is set at 0.1

In the main for-loop, we have 100 iterations which is not a large value. Looping through the training set, for each misclassified point, we use that point to update our weights.

Then we use our model to make a very first prediction:

```
[187] def which_flower(sl, sw, pl, pw):
        return np.sign(sl*w1 + sw*w2 + pl*w3 + pw*w4 + b)

        which_flower(5,3,1,0.5)
```

```
↳ -1.0
```

```
[188] import sklearn.metrics as sm
      y_pred = []
      for i in X_test:
          y_pred.append(which_flower(i[0], i[1], i[2], i[3]))
      y_pred = np.array(y_pred)

      sm.accuracy_score(y_pred, y_test)

0.9333333333333333
```

Afterwards, we use obtained weights to make predictions on the test set, then use the metric `accuracy_score` to compare with the `y_test` labels. Our model is 93.33% accurate on the testing set, which is a great result.

1.2 Using PCA for 2D visualization:

For easier visualization, we apply the PCA technique to reduce our 4-dimensional Iris dataset into a 2-dimensional dataset.

```
# _____ Using PCA _____ #

from sklearn.decomposition import PCA
from sklearn import preprocessing

scaled_data = preprocessing.scale(X)

pca = PCA(n_components=2) # create a PCA object
pca.fit(scaled_data) # do the math
X_pca = pca.transform(scaled_data) # get PCA coordinates for scaled_data

print(X_pca.shape)
print(y.shape)

(100, 2)
(100,)
```

Then, we re-build the Perceptron algorithm with the new 2-dimensional dataset. This time, we have only 2 weights w_1, w_2 and 1 bias b to update.

```

▶ def h(pc1, pc2, w1, w2, b):
    return np.sign(pc1*w1 + pc2*w2 + b)

w1 = 0
w2 = 0
b = 0
alpha = 1

for j in range(1000):
    for i in range(0, len(X_pca)):
        if h(X_pca[i][0], X_pca[i][1], w1, w2, b) != y[i]:
            w1 = w1 + alpha*y[i]*X_pca[i][0]
            w2 = w2 + alpha*y[i]*X_pca[i][1]
            b = b + alpha*y[i]

print(w1)
print(w2)
print(b)

```

```

↳ 2.873257795532451
   3.899221277224042
   2

```

```

[ ] def which_flower(pc1, pc2):
    return np.sign(pc1*w1 + pc2*w2 + b)

which_flower(2,3)

```

```

↳ 1.0

```

```

[ ] #MESH
from matplotlib.colors import ListedColormap
cmap_light = ListedColormap(['#AAAAFF', '#AAFFAA', '#FFAAAA'])
#generate all the points in the plane using np.meshgrid(np.arange(x_min, x_max, super_small))
xx, yy = np.meshgrid(np.linspace(X_pca[:,0].min(), X_pca[:,0].max(), num=100), np.linspace(X_pca[:,1].min(), X_pca[:,1].max(), num=100))
Z = which_flower(xx.ravel(),yy.ravel())
Z = Z.reshape(xx.shape)

plt.figure()
# xx, yy are ALL the points in the plane, Z is the color of each area
# then plot the points using plt.scatter as usual
plt.pcolormesh(xx,yy,Z,cmap=cmap_light)

# Graph points
for i in range(0, len(X_pca)):
    if y[i] == 1:
        plt.scatter(X_pca[i][0], X_pca[i][1], c = 'r')
    elif y[i] == -1:
        plt.scatter(X_pca[i][0], X_pca[i][1], c = 'b')
plt.grid()
plt.xlabel('PC1')
plt.ylabel('PC2')


plt.xlim(xx.min(),xx.max())
plt.ylim(yy.min(),yy.max())
#-----x_min x_max y_min y_max

```

```

↳ (-1.7161812507853758, 2.3299089668338917)

```



We illustrate the final result with mesh color plot using matplotlib.pyplot. It is obvious that the hyperplane do a great job in dividing 2 distinct classes.

Comment on the convergence rate of Perceptron on Iris:

- The convergence rate of the dataset relies on our initialized weights' values, number of iterations and most importantly, the learning rate alpha.
- For this particular dataset, we would say that it converges relatively fast. Given $\alpha=1$, all the weights and bias is initially set at 0, the algorithm gives 93% accurate predictions after just 100 iterations. And as we continue to experiment, our algorithm reaches 97% accurate predictions at 1000 iterations.

For further study, we implement Perceptron with Breast Cancer dataset. But this time, we use the built-in sklearn's perceptron for convenience.

```
[ ] # Breast cancer dataset

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

column = ["Sample code number", "Clump Thickness", "Uniformity of Cell Size", "Uniformity of Cell Shape", "Marginal Adhesion", "Single Epithelial Cell Size",
          "Bare Nuclei", "Bland Chromatin", "Normal Nucleoli", "Mitoses", "Class"]
df = pd.read_csv("/content/gdrive/My Drive/breast-cancer-wisconsin.data", sep=',', header=None, skipinitialspace=True)
df.columns = column

# Save the class information
myclass = df["Class"]

# Set the ID number as index
df.set_index("Sample code number", inplace=True)

# Get rid of class column for statistical analysis and PCA, append it latter
df.drop(["Class"], axis=1, inplace=True)

df.head()
```

	Clump Thickness	Uniformity of Cell Size	Uniformity of Cell Shape	Marginal Adhesion	Single Epithelial Cell Size	Bare Nuclei	Bland Chromatin	Normal Nucleoli	Mitoses
Sample code number									
1000025	5	1	1	1	2	1	3	1	1
1002945	5	4	4	5	7	10	3	2	1
1015425	3	1	1	1	2	2	3	1	1
1016277	6	8	8	1	3	4	3	7	1
1017023	4	1	1	3	2	1	3	1	1

This dataset has 2 classes (benign or malignant) given 9 different attributes ranging from 1 to 10.

We again prepare our X (data) and y (label), divide them into train/test set with the ratio of 90:10.

Then, we import the Perceptron from sklearn.linear_model. The perceptron.coef_ gives all the updated weights.

We use to model to make our predictions on the test set and calculate the accuracy score. 98.57% is a really good score for prediction.

```
[ ] X = df.replace("?", 1).values
    y = myclass.values

    from sklearn.model_selection import train_test_split
    X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.1,random_state=4)
```

```
[ ] from sklearn.linear_model import Perceptron

    perceptron = Perceptron(tol=1e-3, random_state=0)
    perceptron.fit(X_train, y_train)

    # The obtained weights after running the algorithm
    print(perceptron.coef_)
```

```
↳ [[-4. 19.  6.  4. -2. 15.  5.  6. 11.]]
```

```
[ ] y_pred = perceptron.predict(X_test)
    print(y_pred)
```

```
↳ [4 2 2 2 2 2 2 4 4 4 2 2 2 2 4 4 2 2 2 2 4 2 2 2 4 2 2 4 4 2 4 2 2 4 2
    4 4 2 2 4 4 2 2 2 2 2 2 2 2 4 2 2 2 2 4 4 4 2 2 4 4 2 2 2 4 2]
```

```
[ ] sm.accuracy_score(y_pred, y_test)
```

```
↳ 0.9857142857142858
```

II) K Nearest Neighbors.

2.1 Iris dataset

```
[ ] import numpy as np
    import pandas as pd
    import matplotlib.pyplot as plt

    column = ["sepal length", "sepal width", "petal length", "petal width", "class"]
    df = pd.read_csv("/content/gdrive/My Drive/iris.data", sep=',', header=None)
    df.columns = column
    myclass = df["class"]
    df.head(5)
```

```
↳
```

	sepal length	sepal width	petal length	petal width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

```
[ ] X = df.drop(["class"], axis=1).values
    y = myclass.values
```

```
[ ] # splitting the data into training and test sets (80:20)
    from sklearn.model_selection import train_test_split
    X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=4)
```

```
[ ] print(len(X_train))
    print(len(y_train))
```

```
↳ 120
   120
```

KNN algorithm can work on multiple classes, not binary like Perceptron. Therefore, we quickly prepare our data: read it into dataframe with pandas, divide train/test (ratio 80:20).

We first implement to KNN with $k=5$, then report all the score of the dataset by comparing the prediction with the test labels.

```
[ ] from sklearn.neighbors import KNeighborsClassifier
    import sklearn.metrics as sm

    my_knn = KNeighborsClassifier(n_neighbors=5)
    my_knn.fit(X_train, y_train)

    y_predict = my_knn.predict(X_test)

    target_names= ["Iris-setosa", "Iris-versicolor", "Iris-virginica"]
    print(sm.classification_report(y_predict, y_test, target_names=target_names))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	16
Iris-versicolor	0.80	1.00	0.89	4
Iris-virginica	1.00	0.90	0.95	10
accuracy			0.97	30
macro avg	0.93	0.97	0.95	30
weighted avg	0.97	0.97	0.97	30

It is clear that our model with $k=5$ works well in predicting all 3 classes. The overall accuracy is up to 97%, leaving only 3% classification error.

Then we vary k value from 2 to 19, then calculate the accuracy of each KNN.

The accuracy score remains constant when varying k .


```
[ ] for i in range(2,20):
    my_knn = KNeighborsClassifier(n_neighbors=i)
    my_knn.fit(X_train, y_train)

    y_predict = my_knn.predict(X_test)

    print("Accuracy score with k =", i,"is",sm.accuracy_score(y_predict, y_test))
```

```
↳ Accuracy score with k = 2 is 0.9333333333333333
Accuracy score with k = 3 is 0.9666666666666667
Accuracy score with k = 4 is 0.9666666666666667
Accuracy score with k = 5 is 0.9666666666666667
Accuracy score with k = 6 is 0.9666666666666667
Accuracy score with k = 7 is 0.9666666666666667
Accuracy score with k = 8 is 0.9666666666666667
Accuracy score with k = 9 is 0.9666666666666667
Accuracy score with k = 10 is 0.9666666666666667
Accuracy score with k = 11 is 0.9666666666666667
Accuracy score with k = 12 is 0.9666666666666667
Accuracy score with k = 13 is 0.9666666666666667
Accuracy score with k = 14 is 0.9666666666666667
Accuracy score with k = 15 is 0.9666666666666667
Accuracy score with k = 16 is 0.9666666666666667
Accuracy score with k = 17 is 0.9666666666666667
Accuracy score with k = 18 is 0.9666666666666667
Accuracy score with k = 19 is 0.9666666666666667
```

Then, we normalize the dataset and repeat the same process.
The result comes out nearly the same as applying KNN directly.

```
[ ] from sklearn.preprocessing import normalize
    X_normalized = normalize(X)

    # Split train-test set
    X_train,X_test,y_train,y_test = train_test_split(X_normalized,y,test_size=0.2,random_state=4)

    for i in range(2,20):
        my_knn = KNeighborsClassifier(n_neighbors=i)
        my_knn.fit(X_train, y_train)

        y_predict = my_knn.predict(X_test)

        print("Accuracy score with k =", i,"is",sm.accuracy_score(y_predict, y_test))
```

```
↳ Accuracy score with k = 2 is 0.9666666666666667
Accuracy score with k = 3 is 0.9666666666666667
Accuracy score with k = 4 is 0.9666666666666667
Accuracy score with k = 5 is 0.9666666666666667
Accuracy score with k = 6 is 0.9666666666666667
Accuracy score with k = 7 is 0.9666666666666667
Accuracy score with k = 8 is 0.9333333333333333
Accuracy score with k = 9 is 0.9666666666666667
Accuracy score with k = 10 is 0.9333333333333333
Accuracy score with k = 11 is 0.9666666666666667
Accuracy score with k = 12 is 0.9666666666666667
Accuracy score with k = 13 is 0.9666666666666667
Accuracy score with k = 14 is 0.9666666666666667
Accuracy score with k = 15 is 0.9666666666666667
Accuracy score with k = 16 is 0.9666666666666667
Accuracy score with k = 17 is 0.9666666666666667
Accuracy score with k = 18 is 0.9666666666666667
Accuracy score with k = 19 is 0.9666666666666667
```

Afterwards, we apply PCA before implementing the KNN algorithm.

```
from sklearn.decomposition import PCA
from sklearn import preprocessing

scaled_data = preprocessing.scale(X)

pca = PCA(n_components=2) # create a PCA object
pca.fit(scaled_data) # do the math
X_pca = pca.transform(scaled_data) # get PCA coordinates for scaled_data

# Split train-test set
X_train,X_test,y_train,y_test = train_test_split(X_pca,y,test_size=0.2,random_state=4)

for i in range(2,20):
    my_knn = KNeighborsClassifier(n_neighbors=i)
    my_knn.fit(X_train, y_train)

    y_predict = my_knn.predict(X_test)

    print("Accuracy score with k =", i,"is",sm.accuracy_score(y_predict, y_test))
```

Accuracy score with k = 2 is 0.9
Accuracy score with k = 3 is 0.8666666666666667
Accuracy score with k = 4 is 0.9
Accuracy score with k = 5 is 0.9666666666666667
Accuracy score with k = 6 is 0.9666666666666667
Accuracy score with k = 7 is 0.9666666666666667
Accuracy score with k = 8 is 0.9666666666666667
Accuracy score with k = 9 is 0.9333333333333333
Accuracy score with k = 10 is 0.9666666666666667
Accuracy score with k = 11 is 0.9666666666666667
Accuracy score with k = 12 is 0.9666666666666667
Accuracy score with k = 13 is 0.9666666666666667
Accuracy score with k = 14 is 0.9666666666666667
Accuracy score with k = 15 is 0.9666666666666667
Accuracy score with k = 16 is 0.9666666666666667
Accuracy score with k = 17 is 0.9666666666666667
Accuracy score with k = 18 is 1.0
Accuracy score with k = 19 is 0.9666666666666667

The results is worse compared to previous approaches, due to the fact that PCA creates information loss to the original dataset.

Our KNN algorithm is based mainly on the value k chosen. So to improve our model, we propose a method of using k-cross validation to find the average accuracy after multiple implementation on different folds.

```
[ ] from sklearn.model_selection import KFold

avg_accuracy_list = []
kfold = KFold(10, True, 1)

for i in range(2,20):
    accuracy_list = []

    for train_index, test_index in kfold.split(X):
        my_knn = KNeighborsClassifier(n_neighbors=i)
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]
        my_knn.fit(X_train, y_train)
        y_predict = my_knn.predict(X_test)

        accuracy_list.append(sm.accuracy_score(y_predict, y_test))

    avg_accuracy = sum(accuracy_list)/len(accuracy_list)
    avg_accuracy_list.append(avg_accuracy)
    print("Accuracy with k =",i,"is", avg_accuracy)

print("The maximum value of accuracy: ",np.array(avg_accuracy_list).max())
# -----> k = 8,9,14,15
```

```
➡ Accuracy with k = 2 is 0.9466666666666669
Accuracy with k = 3 is 0.9666666666666668
Accuracy with k = 4 is 0.9533333333333335
Accuracy with k = 5 is 0.9600000000000002
Accuracy with k = 6 is 0.9533333333333335
Accuracy with k = 7 is 0.9533333333333335
Accuracy with k = 8 is 0.9733333333333334
Accuracy with k = 9 is 0.9733333333333334
Accuracy with k = 10 is 0.9600000000000002
Accuracy with k = 11 is 0.9666666666666668
Accuracy with k = 12 is 0.9600000000000002
Accuracy with k = 13 is 0.9666666666666668
Accuracy with k = 14 is 0.9733333333333334
Accuracy with k = 15 is 0.9733333333333334
Accuracy with k = 16 is 0.9600000000000002
Accuracy with k = 17 is 0.9666666666666668
Accuracy with k = 18 is 0.9600000000000002
Accuracy with k = 19 is 0.9600000000000002
The maximum value of accuracy: 0.9733333333333334
```

We divide the dataset into 10 folds, record the accuracy score for each fold then averaging them. After that, we compare all the obtained averaged accuracy score. It is evident that $k = 8, 9, 14, 15$ are the best values for k in our KNN algorithm.

Finally, we apply the Leave One Out – a special case of k-cross validation where number of folds = number of data points, leaving only 1 point for validation.

```
[ ] # Leave 1 data point for validation
    from sklearn.model_selection import LeaveOneOut

    accuracy_list = []
    for train_index, test_index in LeaveOneOut().split(X):
        #print("TRAIN:", train_index, "TEST:", test_index)
        my_knn = KNeighborsClassifier(n_neighbors=5)
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]
        my_knn.fit(X_train, y_train)
        y_predict = my_knn.predict(X_test)

        accuracy_list.append(sm.accuracy_score(y_predict, y_test))

    print("Classification Error:", 1-sum(accuracy_list)/len(accuracy_list))
```

➡ Classification Error: 0.033333333333333326

Classification Error is 3.33% as calculated.